3.modelling

Joao Lopes

2024-04-18

Contents

---

1. BUILD A MODEL

[from https://www.tidymodels.org/start/models/]

```r
library("tidyverse")
library("hexbin")
library("tidymodels")
library("dotwhisker")
library("see")
library("performance")
```

1.1. LOOK AT DATA

```r
#loading data and making changes
set.seed(123)
diamonds_data <- diamonds |>
  mutate(
    log_price = log(price),
    log_carat = log(carat),
    fct_cut = factor(cut, ordered = FALSE),
    .keep = "used"
  ) |>
  slice_sample(n = 1000)
```

```r
#looking at data
ggplot(diamonds_data, aes(x = log_carat, y = log_price)) +
  geom_hex() +
  facet_wrap(~fct_cut, ncol = 1)
```

```r
ggplot(diamonds_data, aes(x = log_carat, y = log_price, color = fct_cut)) +
  geom_smooth(method = "lm", formula = "y ~ x")
```

1.2. FIT A MODEL

```r
#functional form
?linear_reg
```

```r
#estimate or train
lm_fit <- linear_reg() |>
  fit(log_price ~ log_carat + fct_cut, data = diamonds_data)
```

```r
#analyse model
lm_fit
summary(lm_fit$fit)
tidy(lm_fit)
```

```r
#plot results
tidy(lm_fit) |>
  dwplot(
    dot_args = list(size = 2, color = "black"),
    whisker_args = list(color = "black"),
    vline = geom_vline(xintercept = 0, colour = "grey50", linetype = 2))
```

```r
#check assumptions
check_model(lm_fit$fit)
```

1.3. USE MODEL TO PREDICT

```r
#create new points
new_points <- expand.grid(
  log_carat = log(2),
  fct_cut = c("Fair", "Very Good", "Ideal")
  )
new_points


#get prediction
mean_pred <- predict(lm_fit, new_data = new_points)
mean_pred


#get confidence interval
conf_int_pred <- predict(lm_fit,
                         new_data = new_points,
                         type = "conf_int")
conf_int_pred

#plot points
plot_data <-
  new_points |>
  bind_cols(mean_pred) |>
  bind_cols(conf_int_pred)

ggplot(plot_data, aes(x = fct_cut)) +
  geom_point(aes(y = exp(.pred))) +
  geom_errorbar(aes(ymin = exp(.pred_lower),
                    ymax = exp(.pred_upper)),
                width = .2) +
  labs(x = "diamond cut", y = "diamond price")
```

---

2. BUILD A MODEL USING WORKFLOW

[from https://www.tidymodels.org/start/recipes/]

```r
library("tidyverse")
library("nycflights13")
library("tidymodels")
library("glmnet")
```

2.1. LOOK AT DATA

```r
#loading data and making changes
flight_data <-
  flights |>
  mutate(
    # Convert the arrival delay to a factor
    arr_delay = ifelse(arr_delay >= 30, "late", "on_time"),
    arr_delay = factor(arr_delay),
    # We will use the date (not date-time) in the recipe below
    date = lubridate::as_date(time_hour)
  ) |>
  # Include the weather data
  inner_join(weather, by = c("origin", "time_hour")) |>
  # Only retain the specific columns we will use
  select(dep_time, flight, origin, dest, air_time, distance,
         carrier, date, arr_delay, time_hour) |>
  # Exclude missing data
  na.omit() |>
  # For creating models, it is better to have qualitative columns
  # encoded as factors (instead of character strings)
  mutate_if(is.character, as.factor)

#looking at data
glimpse(flight_data)

flight_data |>
  count(arr_delay) |>
  mutate(prop = n/sum(n))
```

2.2. SPLIT DATA

```r
set.seed(222)
data_split <- initial_split(flight_data, prop = 3/4)

train_data <- training(data_split)
test_data  <- testing(data_split)
```

2.3. CREATE WORKFLOW

```r
#new recipe
flights_rec <-
  recipe(arr_delay ~ ., data = train_data)
```

```r
#add role to recipe
flights_rec <-
  flights_rec |>
  update_role(flight, time_hour, new_role = "ID")

#add date features
flights_rec <-
  flights_rec |>
  step_date(date, features = c("dow", "month")) |>
  step_holiday(date, holidays = timeDate::listHolidays("US")) |>
  step_rm(date)

#convert categorical variables
flights_rec <-
  flights_rec |>
  step_dummy(all_nominal_predictors())

#remove columns with zero variance
flights_rec <-
  flights_rec |>
  step_zv(all_predictors())

#full recipe
flights_rec <-
  recipe(arr_delay ~ ., data = train_data) |>         #new recipe
  update_role(flight, time_hour, new_role = "ID") |> #add role to recipe
  step_date(date, features = c("dow", "month")) |>    #add date features 1
  step_holiday(date, holidays = timeDate::listHolidays("US")) |> #add date features 2
  step_rm(date) |>                                     #remove variable date
  step_dummy(all_nominal_predictors()) |>             #convert categorical variables
  step_zv(all_predictors())                           #remove columns with zero variance
```

2.4. FIT A MODEL

```r
#build model specification
lr_mod <-
  logistic_reg() |>
  set_engine("glm")
lr_mod

#build workflow
flights_wflow <-
  workflow() |>
  add_model(lr_mod) |>
  add_recipe(flights_rec)
flights_wflow

#fit the model
flights_fit <-
  flights_wflow |>
  fit(data = train_data)

flights_fit |>
```

```
  extract_fit_parsnip() |>
  tidy()
```

2.5. EVALUATE MODEL

```
#predict using augment()
flights_aug <-
  augment(flights_fit, new_data = test_data)
flights_aug |>
  select(arr_delay, time_hour, flight, .pred_class, .pred_on_time)
```

```
#evaluate model using ROC curve
??autoplot
```

```
flights_aug |>
  roc_curve(truth = arr_delay, .pred_late) |>
  autoplot()
```

```
flights_aug |>
  roc_auc(truth = arr_delay, .pred_late)
```

```
#extract results
flights_fit |>
  extract_fit_parsnip() |>
  tidy() |>
  arrange(desc(abs(statistic))) |>
  slice_head(n = 20) |>
  mutate(
    importance = abs(statistic),
    term = reorder(term, importance)
  ) |>
  ggplot(aes(x = importance, y = term)) +
  geom_col()
```

7

3. BUILD VARIOUS MODELS USING WORKFLOW

[from https://www.tidymodels.org/start/case-study/]

```r
library("tidyverse")
library("tidymodels")
library("glmnet")
library("ranger")
library("vip")
```

3.1. LOOK AT DATA

```r
#loading data and making changes
hotels <-
  read_csv("https://tidymodels.org/start/case-study/hotels.csv") |>
  mutate(across(where(is.character), as.factor))

#looking at data
glimpse(hotels)

hotels |>
  count(children) |>
  mutate(prop = n/sum(n))
```

3.2. SPLIT DATA

```r
set.seed(123)
splits      <- initial_split(hotels, strata = children, prop = 0.75)

hotel_other <- training(splits)
hotel_test  <- testing(splits)

hotel_other |>
  count(children) |>
  mutate(prop = n/sum(n))

hotel_test  |>
  count(children) |>
  mutate(prop = n/sum(n))

set.seed(234)
val_set <- validation_split(hotel_other, strata = children, prop = 0.80)
```

3.3. CREATE WORKFLOW AND FIT MODEL 1

```r
#build model
?logistic_reg

lr_mod <-
  logistic_reg(penalty = tune(), mixture = 1) |>
  set_engine("glmnet")
lr_mod
```

```r
#create recipe
holidays <- c("AllSouls", "AshWednesday", "ChristmasEve", "Easter",
              "ChristmasDay", "GoodFriday", "NewYearsDay", "PalmSunday")
lr_recipe <-
  recipe(children ~ ., data = hotel_other) |>
  step_date(arrival_date) |>
  step_holiday(arrival_date, holidays = holidays) |>
  step_rm(arrival_date) |>
  step_dummy(all_nominal_predictors()) |>
  step_zv(all_predictors()) |>
  step_normalize(all_predictors())

#creat workflow
lr_workflow <-
  workflow() |>
  add_model(lr_mod) |>
  add_recipe(lr_recipe)
lr_workflow

#tune and fit model
lr_reg_grid <- tibble(penalty = 10^seq(-4, -2, length.out = 20))
lr_reg_grid

lr_res <-
  lr_workflow |>
  tune_grid(val_set,
            grid = lr_reg_grid,
            control = control_grid(save_pred = TRUE),
            metrics = metric_set(roc_auc))

#select best parameters
lr_res |>
  collect_metrics() |>
  ggplot(aes(x = penalty, y = mean)) +
  geom_point() +
  geom_line() +
  ylab("Area under the ROC Curve") +
  scale_x_log10(labels = scales::label_number())

lr_res |>
  show_best(metric = "roc_auc", n = 5) |>
  arrange(penalty)

lr_best <-
  lr_res |>
  select_best(metric = "roc_auc")
lr_best

#evaluate model using ROC curve
lr_auc <-
  lr_res |>
  collect_predictions(parameters = lr_best) |>
  roc_curve(children, .pred_children) |>
```

```
  mutate(model = "Logistic Regression")
autoplot(lr_auc)
```

3.4. CREATE WORKFLOW AND FIT MODEL 2

```
#build model
?rand_forest
```

```
cores <- parallel::detectCores()

rf_mod <-
  rand_forest(mtry = tune(), min_n = tune(), trees = 1000) |>
  set_engine("ranger", num.threads = cores) |>
  set_mode("classification")
rf_mod

#create recipe
holidays <- c("AllSouls", "AshWednesday", "ChristmasEve", "Easter",
              "ChristmasDay", "GoodFriday", "NewYearsDay", "PalmSunday")
rf_recipe <-
  recipe(children ~ ., data = hotel_other) |>
  step_date(arrival_date) |>
  step_holiday(arrival_date, holidays = holidays) |>
  step_rm(arrival_date)

#creat workflow
rf_workflow <-
  workflow() |>
  add_model(rf_mod) |>
  add_recipe(rf_recipe)
rf_workflow

#tune and fit model
set.seed(345)
rf_res <-
  rf_workflow |>
  tune_grid(val_set,
          grid = 10,
          control = control_grid(save_pred = TRUE),
          metrics = metric_set(roc_auc))

#select best parameters
autoplot(rf_res)

rf_res |>
  show_best(metric = "roc_auc", n = 5)

rf_best <-
  rf_res |>
  select_best(metric = "roc_auc")
rf_best

#evaluate model using ROC curve
```

```
rf_auc <-
  rf_res |>
  collect_predictions(parameters = rf_best) |>
  roc_curve(children, .pred_children) |>
  mutate(model = "Random Forest")
autoplot(rf_auc)
```

3.5. EVALUATE LAST MODEL

```
#compare best models
bind_rows(rf_auc, lr_auc) |>
  ggplot(aes(x = 1 - specificity, y = sensitivity, col = model)) +
  geom_path(lwd = 1.0, alpha = 0.8) +
  geom_abline(lty = 3) +
  coord_equal() +
  theme_bw() +
  theme(legend.position = "bottom")

#build last model
last_rf_mod <-
  rand_forest(mtry = rf_best$mtry, min_n = rf_best$min_n, trees = 1000) |>
  set_engine("ranger", num.threads = cores, importance = "impurity") |>
  set_mode("classification")

#create last workflow
last_rf_workflow <-
  rf_workflow |>
  update_model(last_rf_mod)

#fit last model
set.seed(345)
last_rf_fit <-
  last_rf_workflow |>
  last_fit(split = splits)

#evaluate model using ROC curve
last_rf_fit |>
  collect_predictions() |>
  roc_curve(children, .pred_children) |>
  autoplot()

last_rf_fit |>
  collect_metrics()

#extract results
last_rf_fit |>
  extract_fit_parsnip() |>
  vip(num_features = 20)
```